


Verkle Trees 101


EF Stateless consensus team



Guillaume Ballet



Ignacio Hagopian



Josh Rudolf



Kevaundray Wedderburn

verkle.info



Verkle Trees for Statelessness

💡 Verkle Trees = Vector Commitments + Merkle Trees

Contents

[1 Introduction](#)

[2 FAQ](#)

[3 Dashboard](#)

[Current main tasks](#)

[Upcoming tasks](#)

[Future milestones](#)

[Open questions](#)

[4 Testnets](#)

[5 Resources](#)

[6 Client Implementations](#)

[Updates shared in Verkle Implementers Calls...](#)

[Latest summary here](#)

[Cryptography and other](#)

[Execution Layer](#)

[Consensus Layer](#)

site maintained by [@rudolf6](#), [@ignaciohagopian](#), and [@gballet](#) (ping with any questions/requests)

Motivation

Motivation

- Stateful applications are complex!
- State in systems create many challenges
- State (usually) only grows with time
- For Ethereum it puts some pressure on core values

Motivation

- Stateful applications are complex!
- State in systems create many challenges
- State (usually) only grows with time
- For Ethereum it puts some pressure on core values
- Before being able to validate blocks:
 - Download all the state, which takes time
 - Save it somewhere, a full node requires $\sim(1\text{TiB}+300\text{GiB})$ disk
- Handling state isn't zk-friendly



Motivation

- Let's build a better (stateless) world...


Motivation


- Let's build a better (stateless) world...
- A new node joining the network:
 - Doesn't have to sync all the state.
 - Doesn't require disk storage for the EL client
- zk-friendly L1:
 - We remove complex data-structure (MPT)
 - We remove heavy use of non-zk-friendly hashes (i.e: Keccak)
- Reduce hardware requirements
- Easier to implement a new (stateless) EL client
- Potentially allow increasing the gas limit
- Might trigger the specialization of protocol roles



TL;DR?

- New kid in town: Execution witness
- Contains all the state needed to execute a block
- Contains a (small!) cryptographic proof it's correct
- "All state needed" includes contract code!





Life is complicated...

- Need to introduce a new cryptography stack
- Need to change the state tree(s) data structure
- Need to change gas accounting
- Need to migrate data from MPT to VKT



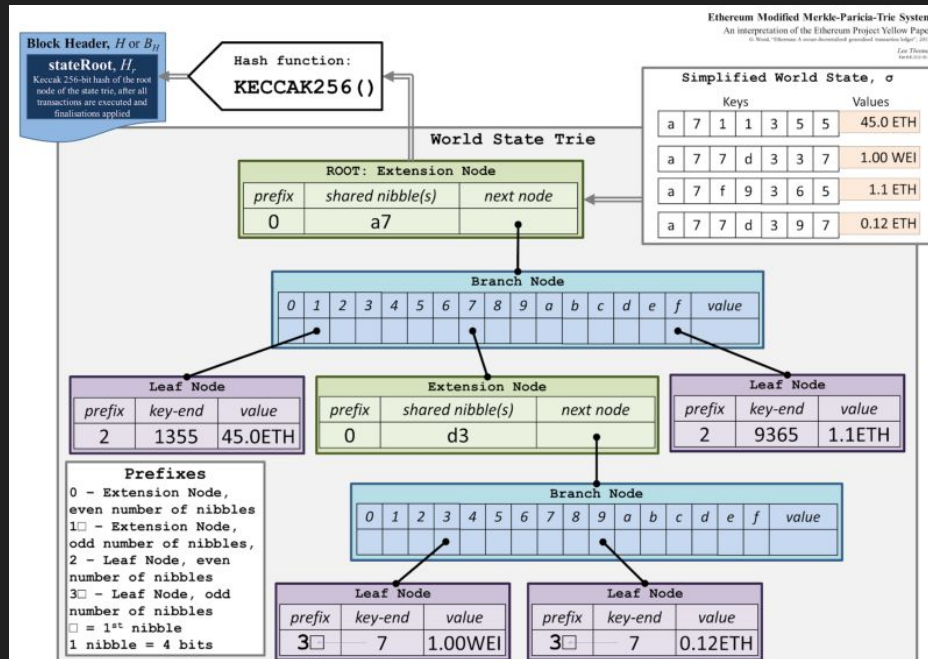
Cryptography

Where the magic happens...

- What makes the *execution witness* proof have a small size:
 - Allows the witness to be transmitted with each block (i.e: each stateless client will need it)
 - Allows keeping the protocol trustless
- Ingredients:
 - *Vector Commitments*
 - *Inner Product Argument*
 - *Multiproofs*


Cryptography used today for the state tree

- Cryptographic hash function: Keccak
- Merkle Patricia Tree:



Cryptography used in Verkle Trees

- Vector commitment:



- Opening:
 $\text{Prove}(V, \text{idx}) = \pi$ (i.e: prove $V[\text{idx}] = \text{res}$)
 $\text{Verify}(C, \text{idx}, \text{res}, \pi) = \text{true/false}$

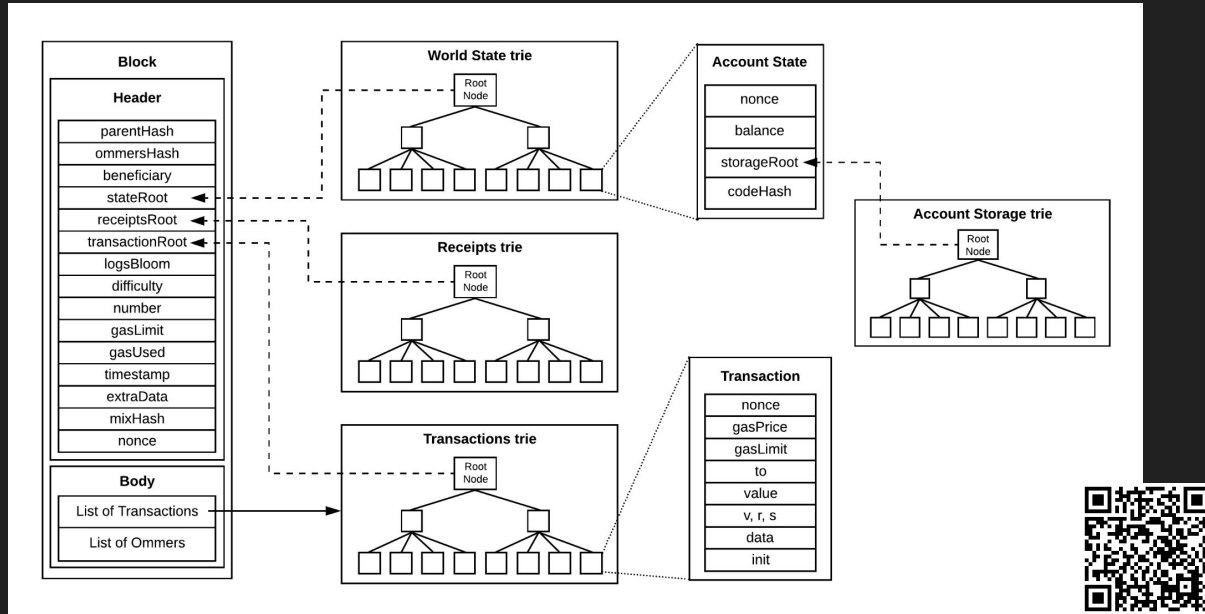
Cryptography used in Verkle Trees

- Group:
 - EC Bandersnatch (Banderwagon, remove cofactor)
 - Scalar field (F_r) = 253 bits
 - Base field (F_p) = 255 bits
 - No pairings (i.e: smaller fields = more efficient)
 - zk-friendly:
 - Base field (F_p) is the Scalar field (F_r) of BLS12-381
 - Doing EC operations in a circuit are native field operations (i.e: not emulating fields)
- Inner Product Argument: single vector opening (no trusted setup!)
- Multiproof: aggregate multiple vector openings in a single vector opening!

Data structure

Data structure

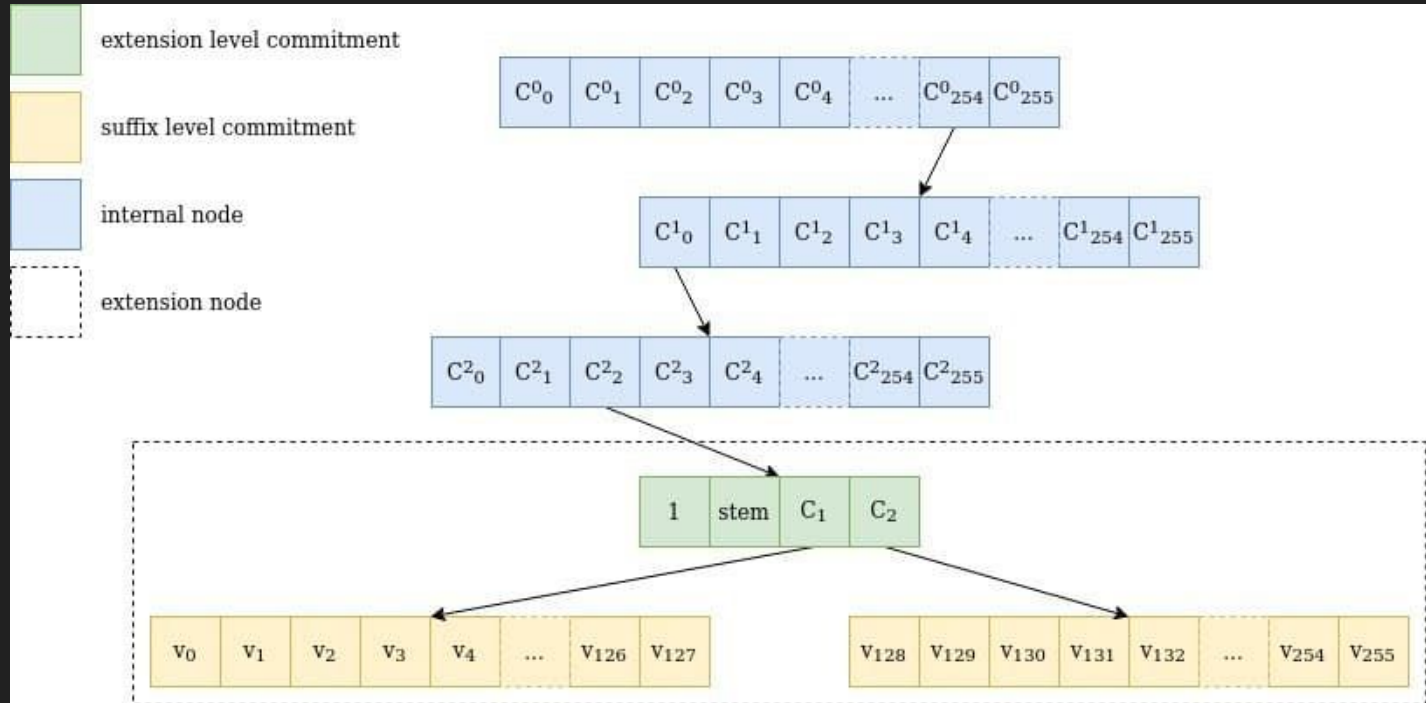
- Merkle Patricia Tree -> Verkle Tree
- Verkle Tree = Vector Commitment + Merkle Tree
- Today:



Data structure (EIP-6800)

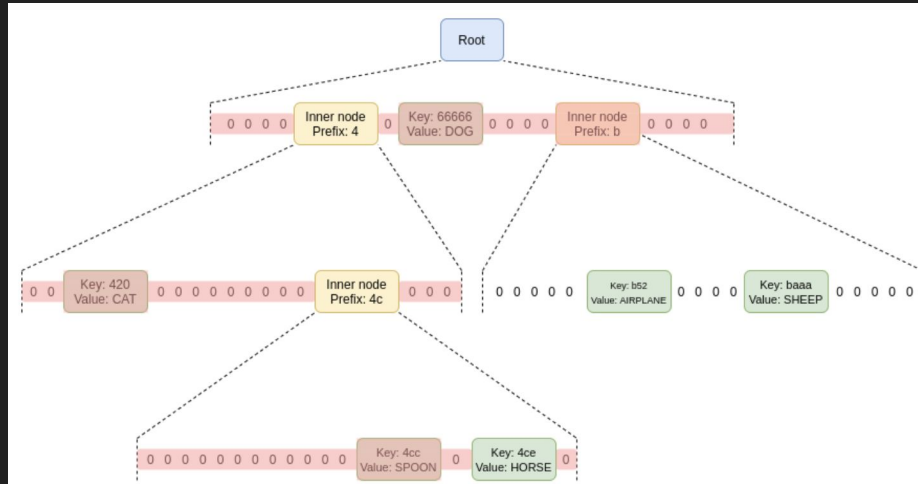
blog.ethereum.org/2021/12/02/verkle-tree-structure

- Future:



$$C_1 = \text{Commit}(v_0^{(\text{lower,modified})}, v_0^{(\text{upper})}, v_1^{(\text{lower,modified})}, v_1^{(\text{upper})}, \dots, v_{127}^{(\text{lower,modified})}, v_{127}^{(\text{upper})})$$
$$C_2 = \text{Commit}(v_{128}^{(\text{lower,modified})}, v_{128}^{(\text{upper})}, v_{129}^{(\text{lower,modified})}, v_{129}^{(\text{upper})}, \dots, v_{255}^{(\text{lower,modified})}, v_{255}^{(\text{upper})})$$

Proving: MPT vs VKT

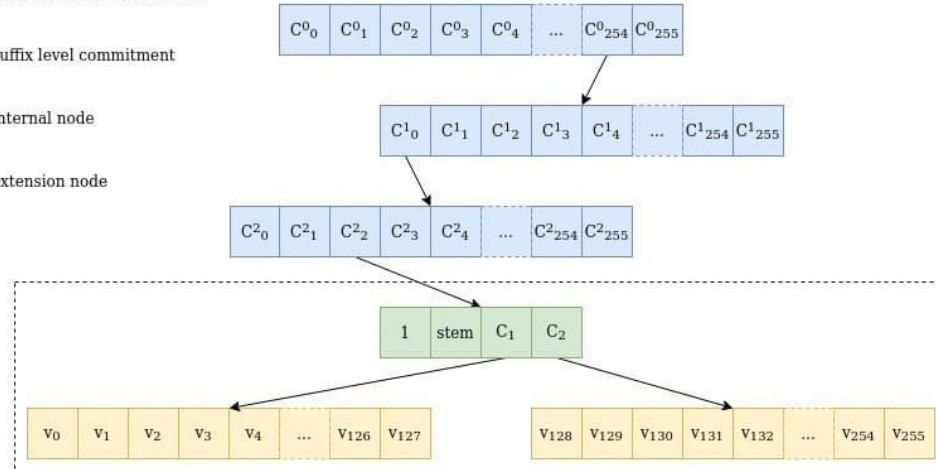


extension level commitment

suffix level commitment

internal node

extension node



<https://vitalik.eth.limo/general/2021/06/18/verkle.html>

How to store information in leaves? (EIP-6800)

- Leaf nodes: store 256 (scalar field) values
- Account header:

Version	Balance	Nonce	CodeHash	CodeSize	...0s...	SS 0	...	SS 63	CodeChunk 0	...	CodeChunk 127
---------	---------	-------	----------	----------	----------	------	-----	-------	-------------	-----	---------------

- Further storage slots:

SS x	SS $x+1$	SS $x+2$...	SS $x+255$
--------	----------	----------	-----	------------

- Let's explain better what *CodeChunk X* means...

How to store information in leaves? (EIP-6800)

- Contracts code is also in the tree!
- Code is “chunked”:
 - A code-chunk is a 32-byte value
 - `Chunk[0]` = How many bytes are a continuation of a PUSHX instruction of prev chunk
 - `Chunk[1:32]` = `Bytecode[0:31]`
- Remember leaves are 256 values of 32 bytes:



How tree keys are calculated? (EIP-6800)

- Stop using Keccak! (not zk-friendly)
- Use a EC based hashing function (i.e: Pedersen hash):

TreeKey(address, treeIndex, subIndex) = Commit(2+256*64, address[0:16], address[16:32], treeIndex[0:16], treeIndex[16:32])[0:31] ++ subIndex

How tree keys are calculated? (EIP-6800)

- Stop using Keccak! (not zk-friendly)
- Use a EC based hashing function (i.e: Pedersen hash):

TreeKey(address, treeIndex, subIndex) = Commit(2+256*64, address[0:16], address[16:32], treeIndex[0:16], treeIndex[16:32])[0:31] ++ subIndex

- Account info: TreeKey(address, 0, [Version | Balance | Nonce | ...])
 - 0 = Version
 - 1 = Balance
 - 2 = Nonce
 - 3 = CodeHash
 - 4 = CodeSize

How tree keys are calculated? (EIP-6800)

- Remember the account header had some storage slots and code chunks:

Version	Balance	Nonce	CodeHash	CodeSize	...0s...	SS 0	...	SS 63	CodeChunk 0	...	CodeChunk 127
---------	---------	-------	----------	----------	----------	------	-----	-------	-------------	-----	---------------

How tree keys are calculated? (EIP-6800)

- Remember the account header had some storage slots and code chunks:

Version	Balance	Nonce	CodeHash	CodeSize	...0s...	SS 0	...	SS 63	CodeChunk 0	...	CodeChunk 127
---------	---------	-------	----------	----------	----------	------	-----	-------	-------------	-----	---------------

- The rest of storage slots:
 - $\text{pos} = \text{MAIN_STORAGE_OFFSET} + \text{storage_key}$
 - TreeKey(address, pos / VERKLE_NODE_WIDTH, pos % VERKLE_NODE_WIDTH)

How tree keys are calculated? (EIP-6800)

- Remember the account header had some storage slots and code chunks:

Version	Balance	Nonce	CodeHash	CodeSize	...0s...	SS 0	...	SS 63	CodeChunk 0	...	CodeChunk 127
---------	---------	-------	----------	----------	----------	------	-----	-------	-------------	-----	---------------

- The rest of storage slots:
 - $\text{pos} = \text{MAIN_STORAGE_OFFSET} + \text{storage_key}$
 - TreeKey(address, pos / VERKLE_NODE_WIDTH, pos % VERKLE_NODE_WIDTH)
- The rest of code chunks:
 - $\text{chunk_id} = \text{CODE_OFFSET} + \text{chunk_id}$
 - TreeKey(address, chunk_id / VERKLE_NODE_WIDTH, chunk_id % VERKLE_NODE_WIDTH)

Gas accounting

Gas accounting (EIP-4762)

- In stateless, IO speed and state size are no longer the primary concern
- Gas should account for increasing the witness size
- Actions that increase witness size
 - Reading state: state needed to be able to execute the block
 - Writing state: provide how the tree looks like in write positions, to be able to update the tree
 - (!!) Executing code: stateless client need the code to execute!

Gas accounting (EIP-4762)

- Accessing a new tree branch -> WITNESS_BRANCH_COST (1900)
- Accessing a new value in a leaf -> WITNESS_CHUNK_COST (200)
- Write triggers updating a branch -> SUBTREE_EDIT_COST (3000)
- Changed value in leaf -> CHUNK_EDIT_COST (500)
- Wrote a leaf node which was empty -> CHUNK_FILL_COST (6200)

Remove the following gas costs:

- Increased gas cost of `CALL` if it is nonzero-value-sending
- EIP-2200 `SSTORE` gas costs except for the `SLOAD_GAS`
- 200 per byte contract code cost

Reduce gas cost:

- `CREATE` to 1000

Gas accounting (EIP-4762)

- Accessing a new tree branch -> WITNESS_BRANCH_COST (1900)
- Accessing a new value in a leaf -> WITNESS_CHUNK_COST (200)
- Write triggers updating a branch -> SUBTREE_EDIT_COST (3000)
- Changed value in leaf -> CHUNK_EDIT_COST (500)
- Wrote a leaf node which was empty -> CHUNK_FILL_COST (6200)
- The above accesses happen when:
 - Particular opcodes are executed (e.g: SSTORE, SLOAD, BALANCE, SELFDESTRUCT, etc)
 - Indirect tree access in transaction executions (e.g: send value FROM to TO) [Not charged!]
 - Indirect tree access in block execution (e.g: withdrawals, block rewards) [Not charged!]
 - Contract code is executed (!!)

Gas accounting (EIP-4762)

- Accessing a new tree branch -> WITNESS_BRANCH_COST (1900)
- Accessing a new value in a leaf -> WITNESS_CHUNK_COST (200)
- Write triggers updating a branch -> SUBTREE_EDIT_COST (3000)
- Changed value in leaf -> CHUNK_EDIT_COST (500)
- Wrote a leaf node which was empty -> CHUNK_FILL_COST (6200)
- Notes:
 - These charges only happen once per key per transaction.
 - Any further storage slot access charges “warm access” cost (100)




State conversion

State conversion



- Need to migrate all the state from MPT -> VKT
- Multiple strategies have been considered
- Currently the “Overlay tree” is the proposed one

Freeze the tree at block $h_{\text{fork}}-1$




- Pre-fork values
- Post-fork values

Start with a fresh root





- Pre-fork values
- Post-fork values

New writes go into the verkle tree



And N leaves are also converted into the verkle tree





- Pre-fork values
- Post-fork values

Accessing data

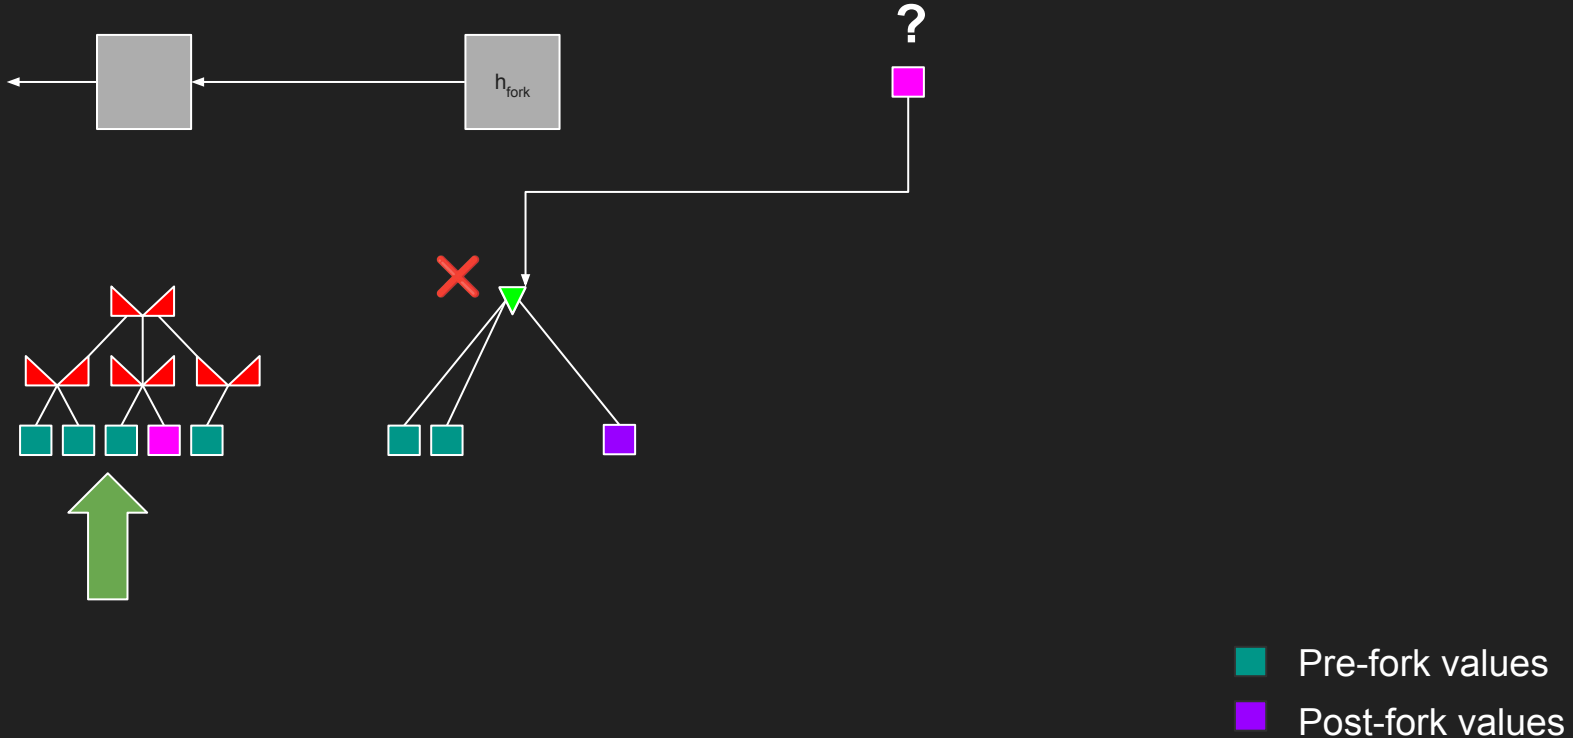


?

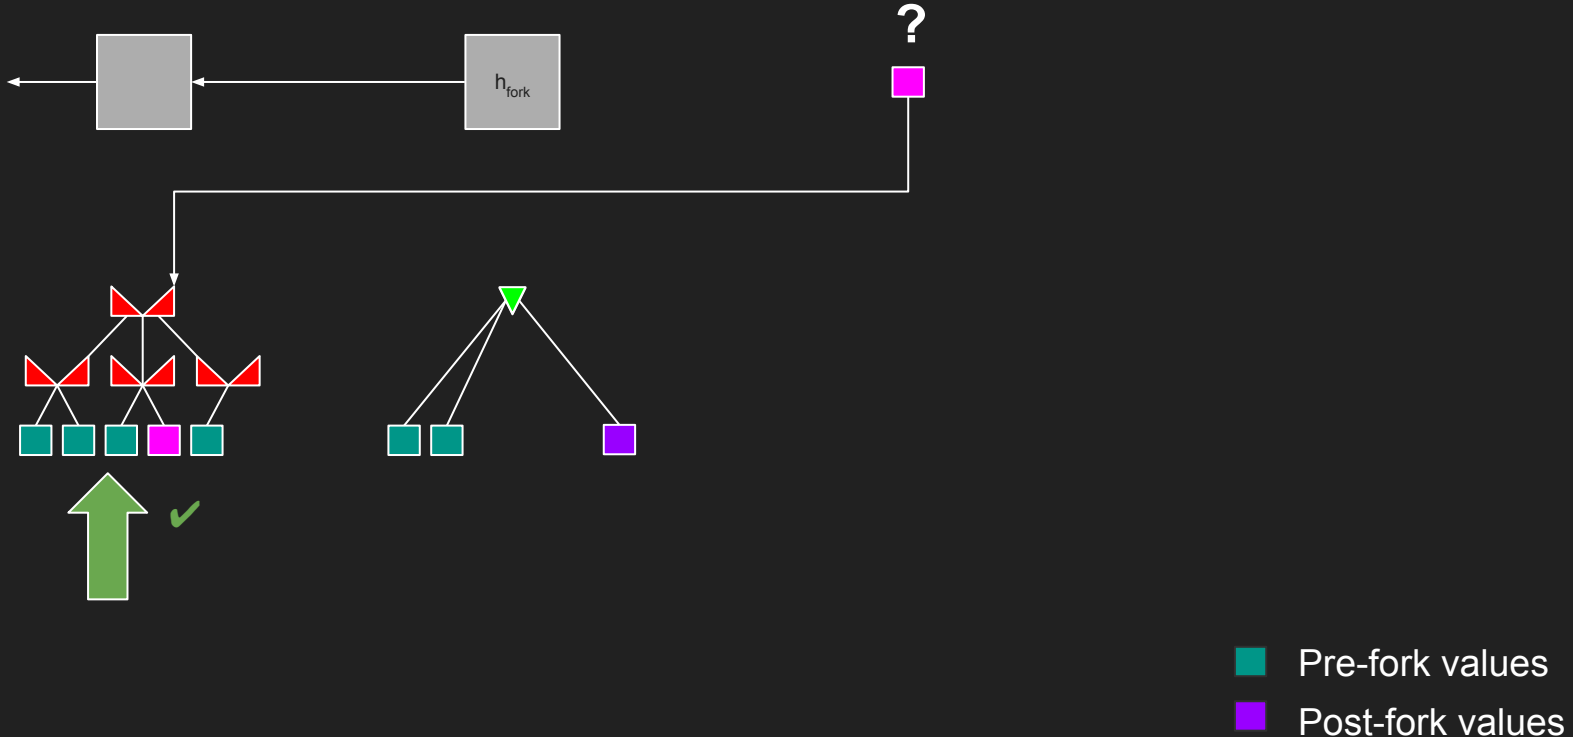


- Pre-fork values
- Post-fork values


Accessing data : first, attempt reading from the verkle tree



Accessing data : if not present, read from the MPT




The same process repeats on the next block




- Pre-fork values
- Post-fork values

Delete internal MPT nodes when the block is finalized




- Pre-fork values
- Post-fork values

Eventually, all the leaves have been converted



■ Pre-fork values
■ Post-fork values

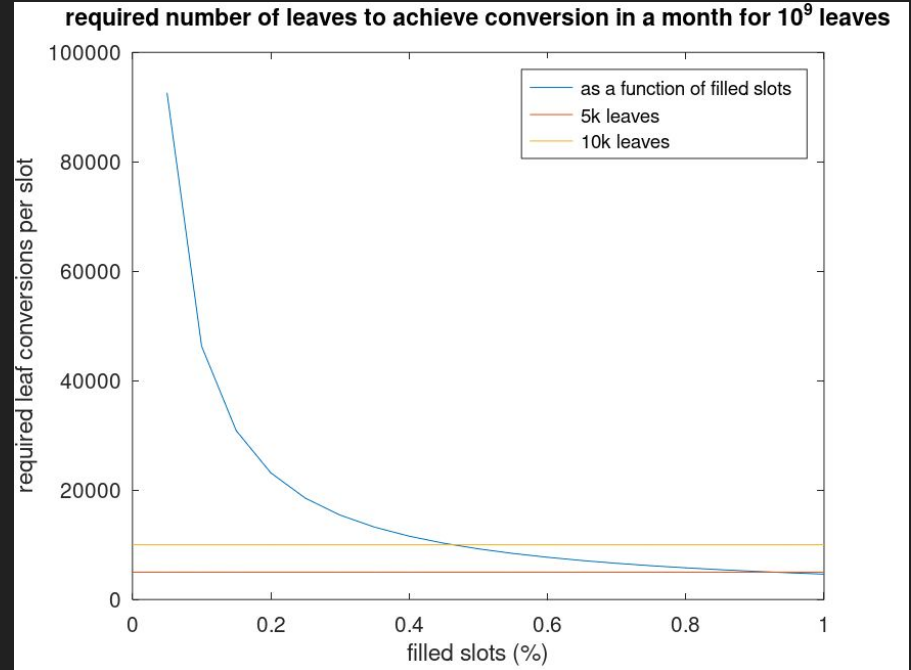
So the MPT data can be forgotten



■ Pre-fork values
■ Post-fork values

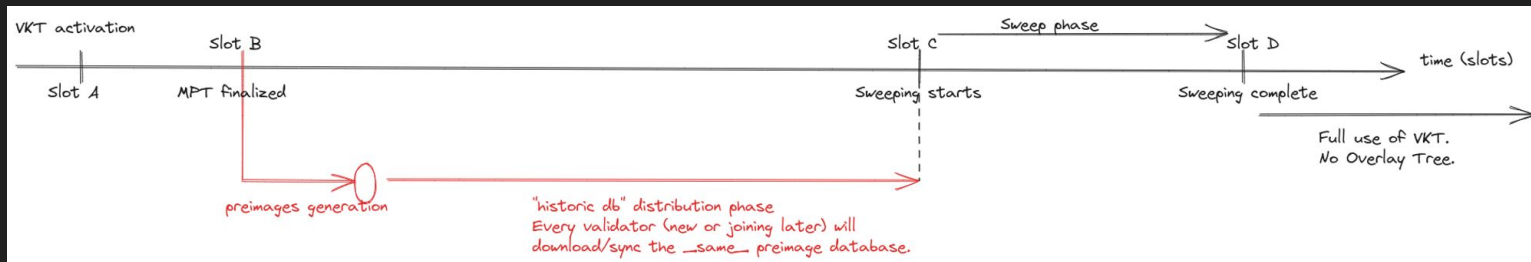
The overlay method : timeline

- $N = 1K \Rightarrow 6$ months
- $N = 5K \Rightarrow 1$ month
- $N = 10K \Rightarrow 15$ days




More challenges

- Re-hashing tree keys requires the pre-images
- Not all EL clients have a design or default flags to store hash pre-images
- How nodes can have this information before the storage transition starts?




hackmd.io/@jsign/vkt-preimage-generation-and-distribution


Verkle sync: view at block n



Verkle sync : view at block n+1



Verkle sync: backfill data



Join the efforts!

- Biweekly calls
- #verkle-trie-migration in Ethereum R&D discord



Thanks!