


By:
Parithosh Jayanthi
DevOps engineer
Ethereum Foundation

Testing and Prototyping Ethereum upgrades





Intro

What was complicated about it?

- >20 client combinations need to be tested & Regressions can sneak in very easily
- Communicating and debugging various client combinations
- Figuring out how to test this in a reliable manner! - We just had single client focussed tests till now
- All future upgrades will inherit some of the complexity - build once, use many
- Competences for ELs and CLs are quite separate

What are devnets?

- A testing mirror of the Ethereum base layer
- Contains EL/CL/Validators, setup in a configuration that we want to test
- Allows devs to deploy forks and changes without affecting mainnet
- Devnets are public, allowing the community to test with us
- A fully in control validator set allows us to design edge case scenarios

What does testing look like today?

- Current upcoming fork (Pectra)
- Upcoming future fork (Verkle)
- Features that are proposed for future forks:
 - ILs
 - EIP-7441 (Whisk)
- Client optimisations:
 - EthereumJS Snap sync testnet
 - Bigboi beaconchain tests for blob/validator limits

Local testing

Hard to test changes quickly!

Fix:

- Earlier we needed full fledged testnets to test changes
- Moving to local devnets enables faster iterations: **enter Kurtosis!**
- Able to work async on features, knowing it'll work with full interop
- Configurable locally: 3s slot times, quick forks, mev workflow
- Scalable: To whatever extent kubernetes/docker allows
- Docs: <https://github.com/kurtosis-tech/ethereum-package>

`"kurtosis run --enclave <name> github.com/kurtosis-tech/ethereum-package
--args-file <filename>"`

Local testing

Configured with YAML:

participants:

- el_type: geth
 cl_type: teku
- el_type: nethermind
 cl_type: prysm

additional_services:

- tx_spammer
- blob_spammer
- dora
- prometheus_grafana

snooper_enabled: true

keymanager_enabled: true

MEV Local testing

mev_type: full

additional_services:

- tx_spammer
- blob_spammer
- custom_flood
- el_forkmon
- dora
- prometheus_grafana

mev_params:

launch_custom_flood: true

mev_relay_image: flashbots/mev-boost-relay:latest

network_params:

seconds_per_slot: 3

How do I prototype?

- Kurtosis works on the concept of “allow everything to be overridden”
- Outside of some network basics, you can change anything in a kurtosis network
- To test protocol changes, we can override the client images
- To test new tools, run an existing network and connect your tool to it
- To test quick forks, we can override the network params

Prototype testing

participants:

- el_type: geth

- el_image: ethpandaops/geth:gballet-transition-post-genesis

- cl_type: lodestar

- cl_image: ethpandaops/lodestar:g11tech-verge

- count: 3

network_params:

- electra_fork_epoch: 1

- genesis_delay: 100

snooper_enabled: true

persistent: true

launch_additional_services: true

additional_services:

- assertoor

- dora

What comes after local testing?

- Remote/public devnets
- We needed a way to set them up easily
- Devnets used to be error prone and time consuming
- Easy drift between setup configs of various testnets due to customizations

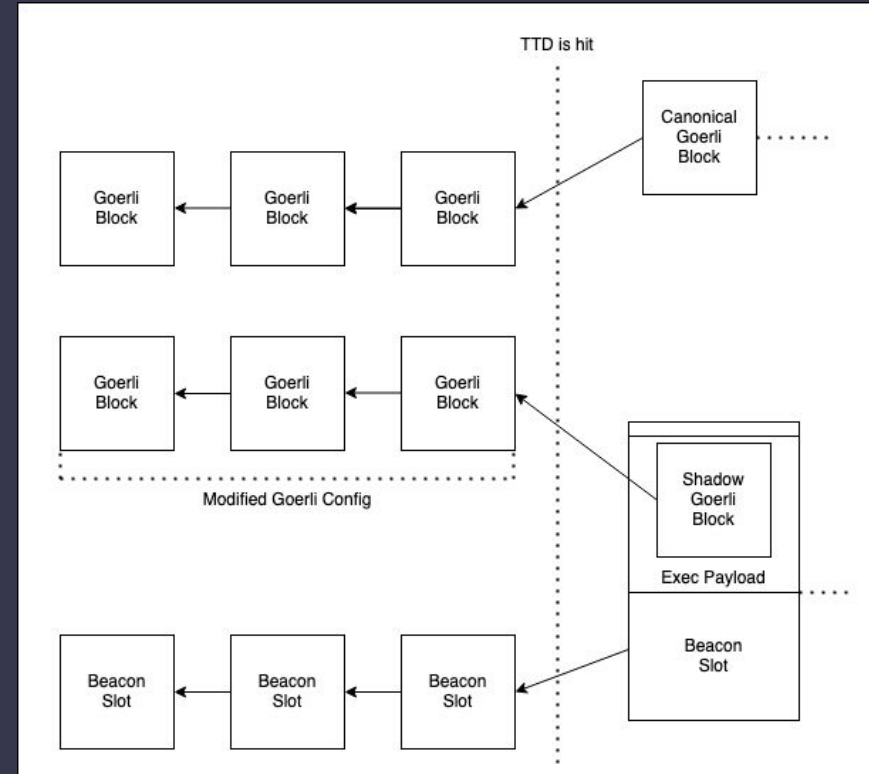
Fix:

- Move barebones logic upstream into role:
<https://github.com/ethpandaops/ansible-collection-general>
- Move generic components into its own tool, e.g: Genesis:
<https://github.com/ethpandaops/ethereum-genesis-generator/>
- Make tooling independent of repo/testnet (With GitOps):
<https://github.com/ethpandaops/ethereum-helm-charts/>
- Generalize setups for all testnets:
<https://github.com/ethpandaops/template-testnet/>

Shadowforks

- Allows us to check compatibility across all clients through the entire lifecycle
- Fresh testnets allow us to check assumptions across client pairs without much overhead
- Shadow forks allow us to stress test the clients with real state and transaction load
- We can invite participants in a controlled manner to take part in the tests
- Acts as release test which triggers real world edge cases, before we recommend the releases to the general public

- Simple in principle:
 - Take a network genesis file
 - Modify file to add a fork timestamp
 - Setup a new beaconchain with validators and same fork timestamp
 - Connect new beaconchain ONLY to ELs with modified config
 - At fork timestamp, the modified ELs and all CLs will shadowfork into a new chain
 - New chain continues to build on top of the canonical chain
- Sprinkle some peering and mempool complexities on top to get a shadowfork





Handy tools to know about

Overview

- Snoop around this website: <https://ethpandaops.io/>
- All codebases can be found here: <https://github.com/ethpandaops/>
- Spend some time reading through:
<https://notes.ethereum.org/@parithosh/testing-overview-doc>
- If you're interested in testing, find a project and start contributing or keep an eye out on `Eth R&D > interop` channel for ideas



Kurtosis (covered earlier)

Template-devnets

- Repo: <https://github.com/ethpandaops/template-testnet/>
- Contains everything you need to configure for any type of testnet
- It uses Terraform to spin up cloud instances and Ansible to deploy the network
- Ansible configs are reliant on these roles:
<https://github.com/ethpandaops/ansible-collection-general>
- Example usage: <https://github.com/ethpandaops/verkle-devnets>
- Useful if you want to run nodes on a larger scale and local testing tools are inadequate

Assertoor

- Repo: <https://github.com/ethpandaops/assertoor>
- Tool to assert network level expectations
- E.g: can a network handle deposits, can it handle every opcode being called, can it handle a reorg
- As its a general testing tool, you can use it for any assertion on any network
- E.g: On verkle networks: was the transition a success
- Similar to hive, Hive -> single node, Assertoor -> Network wide
- Can be run locally, via kurtosis or integrated into a CI:
<https://github.com/ethpandaops/assertoor-github-action>

participants:

- el_type: geth

el_image: ethpandaops/geth:gballet-transition-post-genesis

cl_type: lodestar

cl_image: ethpandaops/lodestar:g11tech-verge

count: 3

network_params:

electra_fork_epoch: 1

network: holesky-shadowfork-verkle

network_sync_base_url: <http://10.10.101.7:9000/snapshots/>

persistent: true

global_node_selectors: {"kubernetes.io/hostname": "lenovo-berlin-02"}

assertoor_params:

image: "ethpandaops/assertoor:verkle-support"

run_stability_check: false

run_block_proposal_check: true

tests:

-

<https://raw.githubusercontent.com/ethpandaops/assertoor-test/master/assertoor-tests/verkle-conversion-check.yaml>

Forky

- Repo: <https://github.com/ethpandaops/forky>
- Ethereum forkchoice visualizer
- Can display the forkchoice of a live node or you can upload your own
- Helps debug forkchoice related issues
- Can be run standalone or via kurtosis
- Forkchoice of mainnet: <https://forky.mainnet.ethpandaops.io/>

Tracoor

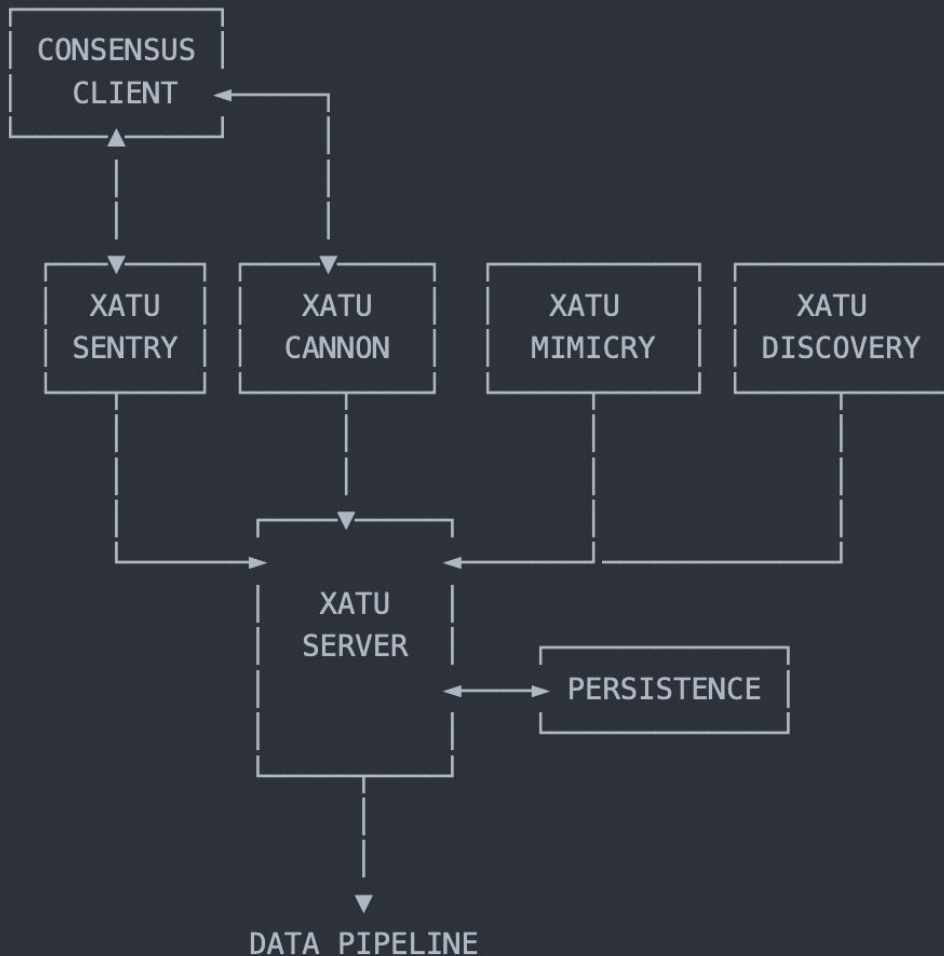
- Repo: <https://github.com/ethpandaops/tracoor>
- Ethereum Trace explorer
- Can display a collection of traces and states of EL and CL blocks/slots
- Helps debug network related issues
- Traces of mainnet: <https://tracoor.mainnet.ethpandaops.io/>

Dora

- Repo: <https://github.com/ethpandaops/dora>
- Lightweight slot explorer for the Ethereum beaconchain
- Extremely extendable, low level access to DB also possible for analysis usage
- Example extendability: <https://dora.verkle-gen-devnet-6.ethpandaops.io/>
- Example standard use: <https://dora.holesky.ethpandaops.io/>

Xatu

- Repo:
 - <https://github.com/ethpandaops/xatu>
 - <https://github.com/ethpandaops/analytics-pipeline>
- Ethereum p2p layer is hard to get visibility about
- Xatu data is then fed into an analysis pipeline to get data we care about
- The visualization is handled by Grafana, but the DB can directly be queried as well
- Data is all open sourced:
 - <https://ethpandaops.io/posts/open-source-xatu-data/>
 - <https://esp.ethereum.foundation/data-challenge-4844>



Xatu Server: central server collecting events from various sentries

Xatu Sentry: client that runs alongside a CL node

Xatu Discovery: client that uses discv4 and discv5 to crawl the network for metadata

Xatu Mimicry: client that collects transaction data from EL p2p layer


Xatu Canon: client that backfills canonical beaconchain data

<https://notes.ethereum.org/@ethpandaops/xatu-overview>

Attestation propagation ▾

29 seconds
28 seconds
27 seconds
26 seconds
25 seconds
24 seconds
23 seconds
22 seconds
21 seconds
20 seconds
19 seconds
18 seconds
17 seconds
16 seconds
15 seconds
14 seconds
13 seconds
12 seconds
11 seconds
10 seconds
9 seconds
8 seconds
7 seconds
6 seconds
5 seconds
4 seconds
3 seconds
2 seconds
1 second
750 milliseconds

Block propagation



Blob sidecar propagation vs number of blobs for slot - all clients



Blob sidecar propagation vs number of blobs for slot - all clients

Blob sidecars	max	p95	average	p50	p05	min
1	7.39 s	2.76 s	767 ms	512 ms	230 ms	118 ms
2	6.60 s	3.00 s	830 ms	553 ms	243 ms	124 ms
3	8.00 s	3.00 s	850 ms	563 ms	259 ms	127 ms
4	8.00 s	2.53 s	859 ms	618 ms	278 ms	130 ms
5	5.00 s	3.48 s	1.38 s	993 ms	409 ms	319 ms
6	7.26 s	2.40 s	852 ms	628 ms	308 ms	162 ms

10/09 16:00 10/09 20:00 10/10 00:00 10/10 04:00 10/10 08:00 10/10 12:00 10/10 16:00 10/10 20:00 10/11 00:00 10/11 04:00 10/11 08:00 10/11 12:00

What next?

Run your own local devnet!

**Get involved in writing tests, best way to
learn how Ethereum works imo**

Thank you!
#TestingThePurge

<https://github.com/ethpandaops/>

@parithosh_j

