

Week 6 Development Track - EL & CL Specs

Guest speaker

- [Hsiao-Wei Wang](#), Ethereum Foundation Consensus R&D team
- [Sam Wilson](#), Ethereum, EIP editor, wallets?

Summary notes

- Edited by [Chloe Zhu](#)
- Online version: <https://ab9jvcjkej.feishu.cn/docx/U9W5dKqIxopyNBxbdf4cXB1InPc>

Consensus Layer Specs

CL specs repo: <https://github.com/ethereum/consensus-specs>

3 purposes of the CL specs repo

- It's a collection of Ethereum core consensus specifications
- It's executable and verifiable
- It's test vector generator

Flow of adding new feature patch

- Implement new features in Pyspec markdown files
- Release new Pyspec with test vector suite
 - Important in the CL R&D process as this process finds basic bugs before implementation
 - Test repo: <https://github.com/ethereum/consensus-spec-tests>
- CL clients implement and test against test vectors

Why use Python?

- Cause it's very readable to devs
- Main principle of the CL specs: Readability & Simplification

How to read the specs?

- Folder structure

- **Specs** folder: incl. all the CL hardforks already happened on the mainnet & WIP research projects
 - Each hardfork folder usually incl. several mark down files, beacon-chain.md would be a good start for each hardfork specs
 - **Features** folder: WIP CL-related research projects
- **SSZ** folder: SSZ (Simple Serialize) containers
 - Info on SSZ: <https://ethereum.org/en/developers/docs/data-structures-and-encoding/ssz/>
 - Devs also use SSZ hash tree root as the digests of consensus objects
- How to read the specs
 - Type and Values definitions: Values are usually defined in the markdown tables, incl. custom types, constants, preset, configuration

Notation

Code snippets appearing in `this style` are to be interpreted as Python 3 code.

Custom types

We define the following Python custom types for type hinting and readability:

Name	SSZ equivalent	Description
Slot	uint64	a slot number
Epoch	uint64	an epoch number
CommitteeIndex	uint64	a committee index at a slot
ValidatorIndex	uint64	a validator registry index

Constants

The following values are (non-configurable) constants used throughout the specification.

Misc

Name	Value
GENESIS_SLOT	Slot(0)
GENESIS_EPOCH	Epoch(0)
FAR_FUTURE_EPOCH	Epoch($2^{64} - 1$)

Preset

Note: The below configuration is bundled as a preset: a bundle of configuration variables which are expected to differ between different modes of operation, e.g. testing, but not generally between different networks. Additional preset configurations can be found in the `configs` directory.

Misc

Name	Value
MAX_COMMITTEES_PER_SLOT	uint64(2^{16}) (= 64)
TARGET_COMMITTEE_SIZE	uint64(2^{21}) (= 128)
MAX_VALIDATORS_PER_COMMITTEE	uint64(2^{11}) (= 2,048)
SHUFFLE_ROUND_COUNT	uint64(90)

Configuration

Note: The default mainnet configuration values are included here for illustrative purposes. Defaults for this more dynamic type of configuration are available with the presets in the `configs` directory. Testnets and other types of chain instances may use a different configuration.

Genesis settings

Name	Value
PER_GENESIS_ACTIVE_VALIDATOR_COUNT	uint64(2^{14}) (= 16,384)
PER_GENESIS_TIME	uint64(1580828800) (Dec 1, 2020, 12pm UTC)

- State transition function
 - Python function can also be written in code block to describe the consensus rule
 - E.g. the assertion in [phase 0/ beacon-chain.md](#)'s state transition function

Beacon chain state transition function

The post-state corresponding to a pre-state `state` and a signed block `signed_block` is defined as `state_transition(state, signed_block)`. State transitions that trigger an unhandled exception (e.g. a failed `assert` or an out-of-range list access) are considered invalid. State transitions that cause a `uint64` overflow or underflow are also considered invalid.

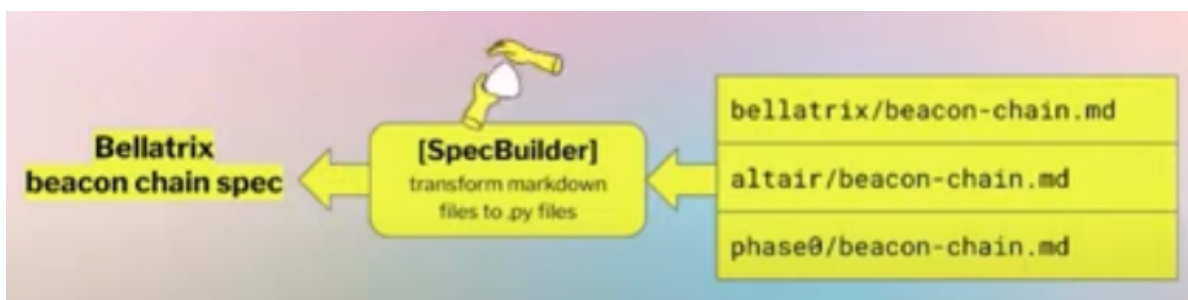
```
def state_transition(state: BeaconState, signed_block: SignedBeaconBlock,   
    block = signed_block.message  
    # Process slots (including those with no blocks) since block  
    process_slots(state, block.slot)  
    # Verify signature  
    if validate_result:  
        assert verify_block_signature(state, signed_block)  
    # Process block  
    process_block(state, block)  
    # Verify state root  
    if validate_result:  
        assert block.state_root == hash_tree_root(state)
```

Useful resources to understand CL

- Vitalik's annotated spec: <https://github.com/ethereum/annotated-spec>
- Ben Edgington's Upgrading Ethereum Book: <https://eth2book.info/capella/>

The elf in setup.py

- Convert the mark down file into Python, and extend the previous hard forks



Demo: How to use Pyspec?

- Installation (Python 3.8+)
 - Download the source code:

```
1 git clone https://github.com/ethereum/consensus-specs.git
```

- Install with Makefile commands

- 1 cd consensus-specs
- 2 make install test && make pyspec

- Run your first pyspec program

```
>> from eth2spec.bellatrix import mainnet as spec

>> hello = b"Hello World"

>> body = spec.BeaconBlockBody(

>>     graffiti=hello + b'\0' * (32 - len(hello))

>> )

>> block = spec.BeaconBlock(body=body)

>> print(block.body.graffiti.decode("utf-8"))

Hello World
```

- Write your first pyspec test case

```
@with_all_phases
@spec_state_test
def test_empty_block_transition(spec, state):
    pre_slot = state.slot
    pre_eth1_votes = len(state.eth1_data_votes)
    pre_mix = spec.get_randao_mix(state, spec.get_current_epoch(state))

    yield 'pre', state

    block = build_empty_block_for_next_slot(spec, state)
    signed_block = state_transition_and_sign_block(spec, state, block)

    yield 'blocks', [signed_block]
    yield 'post', state

    assert len(state.eth1_data_votes) == pre_eth1_votes + 1
    assert spec.get_block_root_at_slot(state, pre_slot) == signed_block.message.parent_root
    assert spec.get_randao_mix(state, spec.get_current_epoch(state)) != pre_mix
```

- Pyspec as the test vector generator

```

@with_all_phases
@spec_state_test
def test_empty_block_transition(spec, state):
    pre_slot = state.slot
    pre_eth1_votes = len(state.eth1_data_votes)
    pre_mix = spec.get_randao_mix(state, spec.get_current_epoch(state))

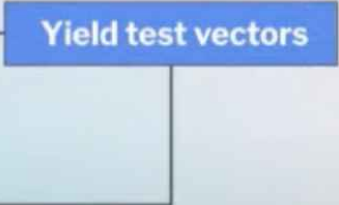
    yield 'pre', state

    block = build_empty_block_for_next_slot(spec, state)
    signed_block = state_transition_and_sign_block(spec, state, block)

    yield 'blocks', [signed_block]
    yield 'post', state

    assert len(state.eth1_data_votes) == pre_eth1_votes + 1
    assert spec.get_block_root_at_slot(state, pre_slot) == signed_block.message.parent_root
    assert spec.get_randao_mix(state, spec.get_current_epoch(state)) != pre_mix

```



The diagram shows a blue box labeled "Yield test vectors" with two arrows pointing to the left. The first arrow points to the line `yield 'pre', state` and the second arrow points to the line `yield 'blocks', [signed_block]` in the code block above.

- Documents
 - Pyspec: <https://github.com/ethereum/consensus-specs/blob/dev/tests/README.md>
 - Test formats: <https://github.com/ethereum/consensus-specs/blob/dev/tests/formats/README.md>

How to contribute to Pyspec?

- Level 1: Look through the specs files to learn about the specs logic & help review it
- Level 2: Help refactor the codebase
- Level 3: Try to hack some new edge test cases
- Level 4: Submit to bug bounty (<https://ethereum.org/en/bug-bounty/>)

Q&A

- How CL specs interact with the EL side? Eg. the state transition function
 - In the fork-choice.md file, the **on_block** function defines whether a block is received. If the assertions are satisfied, then it will call the **state_transition** function from beacon-chain.md
 - Reference file
 - Fork-choice.md: <https://github.com/ethereum/consensus-specs/blob/dev/specs/phase0/fork-choice.md>
 - Beacon-chain.md - State transition function: <https://github.com/ethereum/consensus-specs/blob/dev/specs/phase0/beacon-chain.md#beacon-chain-state-transition-function>
- Any testing for chain reorgs?

- In fork-choice.md, abstract **dataclass** is defined and used to handle reorg and state transition.

```
@dataclass
class Store(object):
    time: uint64
    genesis_time: uint64
    justified_checkpoint: Checkpoint
    finalized_checkpoint: Checkpoint
    unrealized_justified_checkpoint: Checkpoint
    unrealized_finalized_checkpoint: Checkpoint
    proposer_boost_root: Root
    equivocating_indices: Set[ValidatorIndex]
    blocks: Dict[Root, BeaconBlock] = field(default_factory=dict)
    block_states: Dict[Root, BeaconState] = field(default_factory=dict)
    block_timeliness: Dict[Root, boolean] = field(default_factory=dict)
    checkpoint_states: Dict[Checkpoint, BeaconState] = field(default_factory=dict)
    latest_messages: Dict[ValidatorIndex, LatestMessage] = field(default_factory=dict)
    unrealized_justifications: Dict[Root, Checkpoint] = field(default_factory=dict)
```

- Justified checkpoint state is used as the basic stable state
 - **update_checkpoints** function is used to update the justified & finalized checkpoint
 - **get_head** function is used to execute LMD-GHOST fork choice
- In some cases, reorg might happen and **Proposer head and reorg helpers** are gonna help to define the situation and get proposer head
 - Proposer head and reorg helpers
 - is_head_late
 - is_shuffling_stable
 - is_ffg_competitive
 - is_finalization_ok
 - is_proposing_on_time
 - is_head_weak
 - is_parent_strong
 - get_proposer_head

- Reference file: <https://github.com/ethereum/consensus-specs/blob/dev/specs/phase0/fork-choice.md>

- The biggest challenge when developing CL specs?
 - Performance is always a shared challenge. Another one is the conversion between markdown file and python program.

Ethereum Execution Layer Specs (EELS)

EELS repo: <https://github.com/ethereum/execution-specs>

Tldr of the EELS

- From the **Yellow paper** (academic paper with dense math notation) to now **fully executable python implementation of the Ethereum EL**

What's EELS

- It's Python reference implementation of most of a Ethereum client
 - What EELS doesn't do: networking, fork choice, reorgs
 - Other than the above, EELS pretty much builds the entire EL
- Team members
 - Guruprasad Kamath: github.com/gurukamath
 - Peter Davies: <https://github.com/petertdavies>
 - And contributors from all around the world

The Yellow paper and its main problems & benefits

History of the Yellow paper

- Created around 2014 by Gavin Wood
- Creative Commons Attribution Share-Alike (CC-BY-BA) Version 4.0 license, which is a more restrictive one
 - i.e. If you use it, you have to release your stuff under the same license
- The Yellow paper defines the blockchain, fork choice, state, transaction, block, and the rest (incl. Gas, contracts, virtual machine, RLP, modified MPT, precompiles, and EVM instructions)

Main problems

- The math notions in the paper are succinct and difficult to understand
- Audience
 - The Yellow Paper is inaccessible to most programmers, while programmers are the ones who need to understand it the most
 - Core EIP/ EIP authors rarely use this notation
- Untestable
 - The Yellow paper is (mostly) human language
 - No way to test the spec itself or use the spec to fill tests

Key benefits

- Succinct
- Formal

- Algorithm independent

EELS and its main benefits & problems

History of EELS

- Created around May 2021 by Consensys' Quilt team
- Most recently maintained by the EELS team at the Ethereum Foundation
- Creative Commons Zero v1.0 Universal license

Why EELS?

- To escape the frustration of trying to understand the Yellow paper
- No "snapshot" of the current state of Ethereum
- More accessible to programmers with Python
- The same document that specifies Ethereum also fills the automated tests for it

EELS walkthrough

- Directary structure: Tour forks
 - Each fork has its folder and within each folder has the complete copy of Ethereum specs
 - Link: <https://github.com/ethereum/execution-specs/tree/master/src/ethereum>
- Blockchain
 - Regular classes: define what each data structure we need is
 - Python function: define what the behavior is

BlockChain

History and current state of the block chain.

```
72 @dataclass
class BlockChain:
    blocks
    state
    chain_id
```

blocks

```
78     blocks: List[Block]
```

state

```
79     state: State
```

chain_id

```
80     chain_id: U64
```

```
def state_transition(chain: BlockChain, block: Block) -> None:
    167     parent_header = chain.blocks[-1].header
    168     validate_header(block.header, parent_header)
    169     validate_ommers(block.ommers, block.header, chain)
    170     (
    171         gas_used,
    172         transactions_root,
    173         receipt_root,
    174         block_logs_bloom,
    175         state,
    176     ) = apply_body(
    177         chain.state,
    178         get_last_256_block_hashes(chain),
    179         block.header.coinbase,
    180         block.header.number,
    181         block.header.gas_limit,
    182         block.header.timestamp,
    183         block.header.difficulty,
    184         block.transactions,
    185         block.ommers,
    186         chain.chain_id,
    187     )
    188     ensure(gas_used == block.header.gas_used, InvalidBlock)
    189     ensure(transactions_root == block.header.transactions_root, InvalidBlock)
    190     ensure(state_root(state) == block.header.state_root, InvalidBlock)
    191     ensure(receipt_root == block.header.receipt_root, InvalidBlock)
    192     ensure(block_logs_bloom == block.header.bloom, InvalidBlock)
    193
    194     chain.blocks.append(block)
    195     if len(chain.blocks) > 255:
    196         chain.blocks = chain.blocks[-255:]
```

- Forkchoice

```
def validate_proof_of_work(header: Header) -> None:
    300     header_hash = generate_header_hash_for_pow(header)
    303     cache = generate_cache(header.number)
    304     mix_digest, result = hashimoto_light(
    305         header_hash, header.nonce, cache, dataset_size(header.number)
    306     )
    307
    308     ensure(mix_digest == header.mix_digest, InvalidBlock)
    309     ensure(
    310         Uint.from_be_bytes(result) <= (U256_CEIL_VALUE // header.difficulty),
    311         InvalidBlock,
    312     )
```

- EELS doesn't handle reorgs as it assumes only receiving the canonical chain
- State implementation
 - EELS has full state implementation, incl. functions like close_state, begin_transaction, commit_transaction, state_root, storage_root etc.
- Transaction: examples incl.
 - LegacyTransaction: Atomic operation performed on the blockchain
 - AccessListTransaction: The transaction type added in EIP-2930 to support access lists
- Block

Main problems

- Require python knowledge

- Because it's an implementation, it requires spec algorithm choices
- Verbose
- Less accessible to academics

Key benefits

- Much easier to understand
- Audience
 - EELS is accessible to most programmers
 - Core EIPs often use python-style pseudocode. Now it can be actual python
- Maintenance
 - Any programmer can read Python. It's easier for anybody to contribute.
 - Implemented all the way up to Cancun
- Testable
 - Can actually sync the chain (altho very slowly)
 - Pass the ethereum/tests suite
 - Can fill tests for use in production clients
- Cool features
 - Diffs:
 - Tool that generates difference between forks
 - As each hard fork is a complete implementation of the Ethereum protocol. It's hard to tell the difference between forks.
 - Fuzzing
 - Tool that generates inputs randomly and feeded into clients
 - Outputs compared to all other Ethereum clients

Resources

- Adding an EVM instruction to EELS: https://www.youtube.com/watch?v=QlCW_DGSy3s
- The future of EL specs by Peter Davies: <https://archive.devcon.org/archive/watch/6/eels-the-future-of-execution-layer-specifications/?tab=YouTube>

Demo: Add an opcode to EELS

- Video starts at 57:00

Q&A

- How long does it take for EELS to sync the chain?
 - It took roughly 2 weeks to sync the chain. Although it's pretty slow and hard to catch up, it has the capability.
- Regarding the repo, where to find specs post Shanghai hardfork?
 - For Cancun, the repo link: <https://github.com/ethereum/execution-specs/tree/forks/cancun>
- Is the Engine API spec test included under EL specs?
 - API repo: <https://github.com/ethereum/execution-apis>
 - Fork choice: Will add English text, describing how it works.
- Difference between EL and CL specs
 - CL specs approach: start in mark down file, then render into Python
 - EL specs approach: start in Python, and render into HTML
- The biggest challenge to develop EL specs
 - Associated tooling is the hard part, incl. capability to render everything
 - Have to build a whole new rendering system and documentation tool
- Before the EIP/ ERC repo split there was a talk of proposing EIP with EELS. Do we see that in the near future?
 - You can link EELS from EIPs now. But the original vision of moving core EIPs into this repo is kinda dead.